

Fall 2014

University of Southern California

December 22, 2014

Summary

- TA for Machine Learning (50%)
- Two classes
- Meeting...meetings...
- Submitted one paper with Max
- Research project with Fei (collaboration with Yuan)
- Research project with Yan

Scalable inference for overlapping communities




Figure: Batch MCMC vs SG-MCMC vs CGS




Figure: SG-MCMC vs SVI

- Submitted to AISTATS 2015
- Extended version (with 1.8B edges) may be submitted to KDD 2015

Memory Efficient Bayesian Deep Factored Mixed Membership Models




Figure: 3-layer factored LDA

For each document d ,
 for each layer l
 draw (sub)topic distribution π_{dl}
 for each word i
 draw (sub)topic z_{il} for $l \in \{1, \dots, L\}$
 draw word x_i from convex combination of $\{\phi_l\}$

Memory Efficient Bayesian Deep Factored Mixed Membership Models(cont.)

Some benefits:

- Suppose 3-layer model, each layer has L_1, L_2, L_3 components, this can represent $\prod_{i=1}^3 L_i$ mixtures (exponential).
- Less parameters comparing to shallow models.
- Less memory. (becomes very obvious once we have k equals to hundreds of thousand, uses logarithm amount of space comparing to original model)
- Some hierarchical structures

Learning:

- (Stochastic) variational inference, fixed point iteration, L-BFGS
- Or propose better inference

Progress:

Coding, Debugging.... (also need to include deep factored MMSB)

Parallel Stochastic Variational Inference for Bayesian Tensor Factorization

Algorithm 1 Stochastic variational inference

```
1: Initialize  $U, V, W, \alpha, \beta, \gamma$ 
2: while not converge do
3:   Sample a mini-batch of entries,  $\Omega_t$ , from  $X$ 
4:   for  $i \in \Omega_t^i$  do
5:     for  $k=1 \dots K$  do
6:       update  $s_{ik}^u, u_{ik}$  using (8),(9)
7:     end for
8:   end for
9:   for  $j \in \Omega_t^j$  do
10:    for  $k=1 \dots K$  do
11:      update  $s_{jk}^v, v_{jk}$  using (10),(11)
12:    end for
13:  end for
14:  for  $l \in \Omega_t^l$  do
15:    for  $k=1 \dots K$  do
16:      update  $s_{lk}^w, w_{lk}$  using (12),(13)
17:    end for
18:  end for
19:  update  $\alpha, \beta, \gamma$  using (14),(15),(16)
20:  update  $\tau$ 
21: end while
```

Algorithm 2 Parallel Inference

```
1: Initialize sets  $S^u = [1, \dots, I], S^v = [1, \dots, J], S^w = [1, \dots, L]$ 
2: Initialize mini-batch size for each dimension  $N^u = |S^u|/K, N^v = |S^v|/K, N^l = |S^l|/K$ 
3:   where  $K$  is a constant that large than the number of threads  $T$ 
4: while not converge AND not reach max iteration do
5:   for  $t = 1, \dots, T$  parallel do
6:     sample mini-batch of indexes  $I_u^t, I_v^t, I_w^t$  from  $S^u, S^v, S^w$ , the size of mini-batch
       erned by  $N^u, N^v, N^l$ . Then remove  $I_u^t, I_v^t, I_w^t$  from sets  $S^u, S^v, S^w$ 
7:     Form subset of tensor  $\Omega_t$  by  $I_u^t, I_v^t, I_w^t$ 
8:     Update parameters as in Algorithm 1.
9:     Add  $I_u^t, I_v^t, I_w^t$  back into the index sets.
10:   end for
11: end while
```

Extension: distributed settings

Learning from Ordinal Data

- hmm.. this is a hard problem...
- need a simple, elegant solution....
- Working hard on experiments, will tell you more about this.

Plan for winter

- Coding :)
- Debugging :(
- Getting results :)